

ARMA using R

Time series plots

#Band plot

```
library(gplots);  
x=1:1000; y=rnorm(1000, mean=1, sd=1+x/500);  
x11(); bandplot(x, y, main='Band Plot');  
legend('bottomleft' , c('m+/- 2d', 'm+/- d', 'mean'), col=c('magenta', 'blue', 'red'), lwd=c(2,2,2), cex=0.6);
```

#Example

```
ind=labels(AirPassengers);  
x11(); bandplot(ind, AirPassengers, main='Band Plot');  
  
x11(); seasonplot(AirPassengers, col=rainbow(12), year.labels=T);  
x11(); ggseasonplot(AirPassengers, year.labels=T, continuous=T);
```

Prediction (random walk)

Box-Cox transformation

```
library(forecast);  
x11(); plot(AirPassengers, main='Air Passengers'); # original scale  
(lambda=BoxCox.lambda(AirPassengers));  
new=BoxCox(AirPassengers, lambda);  
x11(); plot(new, main='BoxCox Air Passengers'); # scale-changed  
x11(); plot(diff(new), main='Diff-BoxCox Air Passengers');
```

#prediction & evaluation (random walk forecast)

```
library(forecast);  
(rf=rwf(AirPassengers)); # point forecast  
x11(); plot(rf); #prediction interval (80%, 95% upper/lower ends)  
accuracy(rf);
```

white noise

```
set.seed(1234); w=rnorm(300); wt=ts(w);  
x11(); plot(wt, xlab='time', main='white noise'); abline(h=0);  
x11(); acf(wt);  
x11(); pacf(wt);
```

Correcting missing value & outlier

```
# estimation of missing values(multiple imputation by chained equation)
# multiple imputation: Monte Carlo simulation based data merge(random)
x=c(22, 34, NA, 36, 28, 35, 46, 42, 39, 25, 36, 25, 38, NA, 37); t=seq(1:15);
library(mice);
z=mice(matrix(c(x, t), nc=2));
t(complete(z));

# fill the missing values
library(forecast); data(gold);
table(is.na(gold));
for(i in 1: length(gold))
  ifelse(is.na(gold[i]), print(i), next);

# check the number of missing values
# check the location of missing values

tsoutliers(gold);
gold[tsoutliers(gold)$index];
newgold=tsclean(gold);
x11(); tsdisplay(gold);
x11(); tsdisplay(newgold);

# check the location of outliers
# check the outliers
# correct missing values and outliers
# graph before correction
# graph after correction
```

R code for SACF and SPACF

```
library(forecast); library(tseries);
```

AR(1) model

```
x=w=rnorm(500);  
for(t in 2:300) x[t]=0.7*x[t-1]+w[t]; xt=ts(x);  
x11(); par(mfrow=c(1,2)); acf(xt); pacf(xt);
```

ARMA(1, 1) model

```
x=arima.sim(n=500, list(ar=0.7, ma=0.2)); xt=ts(x);  
x11(); par(mfrow=c(1,2)); acf(xt); pacf(xt);
```

ARIMA(1, 1, 1) model

```
x=arima.sim(n=500, list(order=c(1, 1, 1), ar=0.7, ma=0.2)); xt=ts(x);  
x11(); par(mfrow=c(1,2)); acf(xt); pacf(xt);
```

ARIMA(2, 1, 1) model

```
x=arima.sim(n=500, list(order=c(2, 1, 1), ar=c(0.7, 0.2), ma=0.2)); xt=ts(x);  
x11(); par(mfrow=c(1,2)); acf(xt); pacf(xt);
```

Random Walk

```
ddd=arima.sim(list(order=c(0, 1, 0)), n=500); x11(); plot(ddd);  
x11(); par(mfrow=c(1,2)); acf(ddd); pacf(ddd);
```

Unit Root Tests

```
library(tseries);  
### stationary test: unit root test #####  
kpss.test(AirPassengers); # H0: stationary  
# check stationarity after deleting the trend  
x11();plot(AirPassengers); kpss.test(AirPassengers, 'Trend');  
pp.test(AirPassengers); # H0: non-stationary(unit-root)  
adf.test(AirPassengers); # H0: non-stationary(unit-root)  
#####  
x=rnorm(1000); y=diffinv(x); # x has no unit-root, but y contains a unit-root.  
adf.test(x, k=3); pp.test(x); kpss.test(x);  
##### more advanced approach  
y01=read.delim("mraw.txt", header=T);  
m1w1=lm(wti~-1+wti1+snp, data=y01);  
re1=residuals(m1w1);  
library(urca); ut1=ur.df(re1, type= "none", selectlags="AIC"); # lag selection based on AIC criteria  
summary(ut1); # if z.lag.1 is not significant(coefficient of lag1 y=0) , there is unit root  
?ur.df  
#####  
library(tseries); adf.test(re1);  
adf.test(re1, k=2);  
#####  
data(Raotbl3)  
attach(Raotbl3)  
lc.df <- ur.df(y=lc, lags=3, type='trend') # specify the lag  
summary(lc.df)
```

Autocorrelation Tests

- Runs test, Portmanteau test: Box test (Box-Pierce, Ljung-Box)

```
### Runs test
```

```
library(tseries);  
x=rnorm(50); x1=factor(ifelse(x>=median(x), 1, 0));  
runs.test(x1, a='less');
```

```
### Portmanteau test: Box test (Box-Pierce, Ljung-Box)
```

```
Box.test(x, t=c('B')); # Box-Pierce test
```

```
Box.test(x, t=c('L')) # Ljung-Box test
```

```
### arima.sim & arima #####
```

```
y=arima.sim(list(order=c(1,0,0), ar=0.4), n=300); #AR(1) simulation
```

```
x11(); ts.plot(y);
```

```
a=arima(y,order=c(1,0,0)); # AR(1) estimation
```

```
?arima # see method
```

```
forecast(a, h=20);
```

```
predict(a, n.ahead=20);
```

```
Box.test(a$residuals, lag=2);
```

```
Box.test(a$residuals, lag=2, type="Ljung-Box")
```

Co-integration Test & Granger Causality

```
### Phillips-Ouliaris Co-integration test
```

```
x=diffinv(rnorm(1000)); y=2-3*x+rnorm(x, sd=5);  
z=ts(cbind(x,y));      # x and y are co-integrated  
x11(); plot(z);  
po.test(z); # null: no co-integration
```

```
### Johansen Co-integration test: useful for vector AR
```

```
data(denmark) ;  
sjd <- denmark[, c("LRM", "LRY", "IBO", "IDE")] ;  
head(sjd);
```

```
sjd.vecm <- ca.jo(sjd, ecdet = "const", type="eigen", K=2, spec="longrun", season=4) ;  
summary(sjd.vecm);  
?ca.jo
```

```
##### Granger Causality test
```

```
library(lmtest); data(ChickEgg);  
grangertest(chicken~egg, order=3, data=ChickEgg); # egg granger-caused chicken  
grangertest(egg~chicken, order=3, data=ChickEgg); # chicken did not granger-cause egg
```

```
# alternative way to give same result
```

```
grangertest(ChickEgg, order=3);  
grangertest(ChickEgg[,1], ChickEgg[,2], order=3);
```


AR model estimation

```
library(forecast); library(tseries);  
##### Example 1 (no trend) #####  
dd1=c(1342, 1442, 1252, 1343, 1425, 1362, 1456, 1272, 1243, 1359, 1412, 1253, 1201, 1478, 1322,  
1406, 1254, 1289, 1497, 1208);  
d1=ts(dd1, start=c(2016, 1), frequency=4);  
  
# stationary test (unit-root test)  
ddt=d1; kpss.test(ddt); kpss.test(ddt, 'Trend');  
x11(); tsdisplay(ddt, main='time series display');  
  
# ar w/ Yule-walker estimation  
(ar1=ar(ddt, m=c('yule-walker'))); # AR(2) fitting  
(f1=forecast(ar1));  
accuracy(f1);  
  
x11(); plot(f1, xlab='time', ylab='series'); abline(h=mean(ddt));  
grid(); title('Wn Wn estimation: yule-walker');  
Box.test(ar1$resid, type='Ljung-Box');  
temp1=window(ar1$resid, start=c(2016, 3));  
jarque.bera.test(temp1);
```

AR model

```
# ar w/ OLS estimation
```

```
(ar3=ar(ddt, m=c('ols'))); # AR(8) fitting
```

```
(f3=forecast(ar3));
```

```
accuracy(f3);
```

```
x11(); plot(f3, xlab='time', ylab='series'); abline(h=mean(ddt)); grid(); title('\n\n estimation: ols');
```

```
Box.test(ar3$resid, type='Ljung-Box');
```

```
temp3=window(ar3$resid, start=c(2018, 1));
```

```
jarque.bera.test(temp3);
```

```
# ar w/ MLE estimation
```

```
(ar4=ar(ddt, m=c('mle'))); # AR(8) fitting
```

```
(f4=forecast(ar4));
```

```
accuracy(f4);
```

```
x11(); plot(f4, xlab='time', ylab='series'); abline(h=mean(ddt)); grid(); title('\n\n estimation: mle');
```

```
Box.test(ar4$resid, type='Ljung-Box');
```

```
temp4=window(ar4$resid, start=c(2018, 1));
```

```
jarque.bera.test(temp4);
```

```
# fpeaut: AR model optimal lag selection
```

```
library(timsac);
```

```
fpeaut(ddt)$order;
```

```
fpeaut(ddt)$best.ar; # AR(2)
```

ARMA model

```
library(tseries); # arma function
##### Example 1 #####
dam=c(1142, 1242, 1252, 1343, 1225, 1562, 1365, 1572, 1343, 1459, 1412, 1453, 1401, 1478, 1322,
1606, 1554, 1589, 1597, 1408);
damt=ts(dam, start=c(2016,1), frequency=4); kpss.test(damt); kpss.test(damt, 'Trend');
x11(); tsdisplay(damt, main='time series display');

#ARMA(2, 0) model
arm2=arma(damt, order=c(2,0)); summary(arm2); fitted(arm2);
x11(); plot(damt, type='b', main='ARMA (2, 0)'); lines(fitted(arm2), col='red', lty=6, lwd=2); grid();

#ARMA(0, 1) model
arm3=arma(damt, order=c(0,1)); summary(arm3); fitted(arm3);
x11(); plot(damt, type='b', main='ARMA (0, 1)'); lines(fitted(arm3), col='red', lty=6, lwd=2); grid();

#ARMA(1, 1) model
arm11=arma(damt, order=c(1,1)); summary(arm11); fitted(arm11);
x11(); plot(damt, type='b', main='ARMA (0, 2)'); lines(fitted(arm11), col='red', lty=6, lwd=2); grid();

# autoarmafit: ARMA model lag selection
library(timsac);
(aa=autoarmafit(damt)); # Best ARMA model search
ar=aa$model[[1]]$arcoef; ma=aa$model[[1]]$macoef; va=aa$model[[1]]$v;
x11(); prdctr(damt, r=5, s=21, h=4, arcoef=ar, macoef=ma, v=va); title('\n\n Auto ARMA');
```

ARIMA model

Example 1

```
ddd=arima.sim(list(order=c(0, 1, 0)), n=500);  
auto.arima(ddd, allowdrift=F); # optimal model: ARIMA(0, 1, 0)  
am=arima(ddd, order=c(0, 1, 0)); x11(); tsdiag(am);
```

auto.arima: ARIMA model lag selection

Example 2 (no trend)

```
dd1=c(1342, 1442, 1252, 1343, 1425, 1362, 1456, 1272, 1243, 1359, 1412, 1253, 1201, 1478, 1322,  
1406, 1254, 1289, 1497, 1208);  
d1=ts(dd1, start=c(2016, 1), frequency=4);  
am1=auto.arima(d1, allowdrift=F); # optimal model: ARIMA(0, 0, 1)  
x11(); tsdiag(am1);  
fitted(am1); forecast(am1);  
x11(); plot(forecast(am1)); lines(fitted(am1), col='red', lty=2, lwd=2);
```

Example 3 (w/ drift)

```
dd2=c(1142, 1242, 1452, 1543, 1225, 1362, 1556, 1672, 1343, 1459, 1662, 1753, 1421, 1558, 1772,  
1846, 1554, 1649, 1877, 1948);  
d2=ts(dd2, start=c(2016, 1), frequency=4);  
am2=auto.arima(d2); # optimal model: ARIMA(0, 0, 0)(1, 1, 0)[4] with drift  
x11(); tsdiag(am2);  
fitted(am2); forecast(am2);  
x11(); plot(forecast(am2)); lines(fitted(am2), col='red', lty=2, lwd=2);
```

AR(p)I(d)MA(q) model

```
library(forecast); library(tseries);
```

```
### AR(ar_lag1, no_intercept)
```

```
y=arima.sim(list(order=c(1,0,0), ar=0.7), n=300); #AR(1) simulation
```

```
aa=arima(y,order=c(1,0,0))
```

```
# AR estimation: fixed=c(ar1, intercept)
```

```
a=arima(y,order=c(1,0,0), fixed=c(NA, 0), transform.pars=F); #NA: value to estimate, 0: deleted (forced)
```

```
### AR(ar_lag2, ar_lag3)
```

```
y0=arima.sim(list(order=c(3,0,0), ar=c(0, 0.7, 0.2)), n=300);
```

```
# AR estimation: fixed=c(ar1, ar2, ar3, intercept)
```

```
a0=arima(y0,order=c(3,0,0), fixed=c(0, NA, NA, NA), transform.pars=F);
```

```
### ARMA(ar_lag3, ma_lag1, ma_lag2)
```

```
y1=arima.sim(list(order=c(3,0,2), ar=c(0, 0, 0.2), ma=c(0.4, 0.2)), n=300);
```

```
# AR estimation: fixed=c(ar1, ar2, ar3, ma1, ma2, intercept)
```

```
a1=arima(y1, order=c(3,0,2), fixed=c(0, 0, NA, NA, NA, NA), transform.pars=F);
```

```
### ARIMA(ar_lag3, ma_lag1, ma_lag2): If differenced, there is no intercept
```

```
y2=arima.sim(list(order=c(3,1,2), ar=c(0, 0, 0.2), ma=c(0.4, 0.2)), n=300);
```

```
# AR estimation: fixed=c(ar1, ar2, ar3, ma1, ma2)
```

```
a2=arima(y2, order=c(3,1,2), fixed=c(0, 0, NA, NA, NA), transform.pars=F);
```

SARIMA model

```
library(forecast);  
##### Example 1 #####  
x11(); plot(decompose(USAccDeaths));  
x11(); tsdisplay(USAccDeaths, main='US Accidental Deaths / Month');  
findfrequency(USAccDeaths);  
am0=auto.arima(USAccDeaths); summary(am0);  
x11(); tsdiag(am0);  
fitted(am0); forecast(am0, h=30);  
  
##### Example 2 (w/ trend) #####  
dd2=c(1142, 1242, 1452, 1543, 1225, 1362, 1556, 1672, 1343, 1459, 1662, 1753, 1421, 1558, 1772, 1846,  
1554, 1649, 1877, 1948);  
d2=ts(dd2, start=c(2016, 1), frequency=4);  
kpss.test(d2); kpss.test(d2, 'Trend');  
findfrequency(d2);  
  
am=auto.arima(d2); summary(am); # optimal model: ARIMA(0, 0, 0)(1, 1, 0)[4] with drift  
x11(); tsdiag(am);  
fitted(am); forecast(am);  
x11(); plot(forecast(am)); lines(fitted(am), col='red', lty=2, lwd=2);
```

Multiplicative Seasonal ARIMA model

- Co2 data: Monthly Carbon Dioxide Levels at Alert, NWT, Canada

```
data(co2);  
x11(); # 1959/1~1997/12  
win.graph(width=4.875, height=3, pointsize=8);  
plot(co2, ylab='CO2');  
  
# plot with month symbol for subset data  
x11();  
plot(window(co2, start=c(1995,1)), ylab='CO2');  
Month=c('Ja','F','M','A','My','Jn','J','Au','S','O','N','D');  
points(window(co2, start=c(1995,1)), pch=Month);  
  
#sample ACF of co2  
x11(); pacf(as.vector(co2), lag.max=36);  
  
# 1st order difference of co2  
x11(); plot(diff(co2), ylab='1st order difference of co2', xlab='time');  
x11(); pacf(as.vector(diff(co2)), lag.max=36);
```

Multiplicative Seasonal ARIMA model

- Co2 data: Monthly Carbon Dioxide Levels at Alert, NWT, Canada

```
# 1st and seasonal difference of co2
```

```
x11(); plot(diff(diff(co2), lag=12), ylab='1st and seasonal difference of co2', xlab='time');  
x11(); pacf(as.vector(diff(diff(co2), lag=12)), lag.max=36, ci.type='ma');
```

```
# estimation of ARIMA(1,1,1)(1,1,1)12
```

```
auto.arima(co2); # Best model fitting
```

```
M1.co2=arima(co2, order=c(1,1,1), seasonal=list(order=c(1,1,2), period=12)); M1.co2;  
f1=forecast(M1.co2, h=30);  
accuracy(f1);
```

```
# residual analysis
```

```
x11(); plot(window(rstandard(M1.co2), start=c(1995,2)), ylab='standardized residuals', type='o');  
abline(h=0);
```

```
x11(); acf(as.vector(window(rstandard(M1.co2), start=c(1995,2))), lag.max=36);
```

```
x11(); win.graph(width=3, height=3, pointsize=8);  
hist(window(rstandard(M1.co2), start=c(1995, 2)), xlab= 'standardized residuals');
```

```
x11(); win.graph(width=3, height=3, pointsize=8);  
qqnorm(window(rstandard(M1.co2), start=c(1995,2)));  
qqline(window(rstandard(M1.co2), start=c(1995,2)));
```


ARIMAX

```
library(forecast);  
# simulate random data  
x <- ts(rnorm(120,0,3) + 1:120 + 20*sin(2*pi*(1:120)/12), frequency=12);  
temp = rnorm(length(x), 20, 30); # exogenous variable
```

```
# build the ARMAX model  
model1 = auto.arima(x, xreg = temp);
```

```
# model summary  
summary(model1);
```


```
# new predictors: 10 data points  
temp.reg = rnorm(10, 20, 30);  
n=length(temp.reg);  
x11(); plot(1:n, temp.reg, type='l');
```

```
# forecasting  
forecast1 = forecast(model1, xreg = temp.reg);
```



```
# visualize  
x11(); plot(forecast1);
```

```
# model info  
summary(forecast1);
```

Forecasts from Regression with ARIMA(0,0,0)(2,1,1)[12] errors



```
library(TSA);  
data(airmiles);  
x11(); plot(log(airmiles),ylab='Log(airmiles)',xlab='Year', main="") ;  
x11(); plot(acf(diff(diff(window(log(airmiles), end=c(2001,8)), 12)), lag.max=48,main=""));
```



ARIMAX

```
# in-sample & out-of-sample
airmiles_in=airmiles[1:100,]; # 1996.1~2004.4
reg=c(rep(0,11),1,1,rep(0,87)) # dummy regress: 100 points
airmiles_out=airmiles[101:113,]; # 2004.5~2005.5
```

```
# build the ARMAX model using in-sample
model2=auto.arima(log(airmiles_in), xreg=reg);
```

```
# model summary
summary(model2);
```

```
# new predictors: new dummy regress: 13 points
new.reg=c(rep(0, 7), 1, rep(0, 5));
```

```
# forecasting
forecast2 = forecast(model2, xreg = new.reg);
```

```
# prediction accuracy
x11(); plot(forecast2);
pred=exp(forecast2$mean);
diff=(airmiles_out-pred)^2;
mse=mean(diff);
rmse=sqrt(mse);
rmse/mean(airmiles);
```

Forecasts from Regression with ARIMA(1,1,1) errors

