

Ch 3: Classifiers using scikit-learn

□ topics

- popular algorithms for classification such as logistic regression, support vector machines, and decision trees
- using the scikit-learn machine learning library
- discussions on classifiers with linear and nonlinear decision boundaries

scikit-learn API

- ❑ combines a user-friendly and consistent interface with a highly optimized implementation of several classification algorithms.
- ❑ scikit-learn library offers
 - a large variety of learning algorithms
 - many convenient functions to preprocess data and to fine-tune and evaluate models.

scikit-learn API

```
from sklearn import datasets
import numpy as np


iris = datasets.load_iris()
X = iris.data[:, [2, 3]]
y = iris.target

print('Class labels:', np.unique(y))
```

Class labels: [0 1 2]

scikit-learn API

```
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.3, random_state=1, stratify=y)
```



scikit-learn API

Standardizing the features:

```
from sklearn.preprocessing import StandardScaler  
  
sc = StandardScaler()  
sc.fit(X_train)  
X_train_std = sc.transform(X_train)  
X_test_std = sc.transform(X_test)
```

Training a perceptron via scikit-learn

```
from sklearn.linear_model import Perceptron  
  
ppn = Perceptron(eta0=0.1, random_state=1)  
ppn.fit(X_train_std, y_train)
```

```
y_pred = ppn.predict(X_test_std)  
print('Misclassified examples: %d' % (y_test != y_pred).sum())
```

Misclassified examples: 1

```
from sklearn.metrics import accuracy_score  
  
print('Accuracy: %.3f' % accuracy_score(y_test, y_pred))
```

Accuracy: 0.978

```
print('Accuracy: %.3f' % ppn.score(X_test_std, y_test))
```

Accuracy: 0.978

```

from matplotlib.colors import ListedColormap
import matplotlib.pyplot as plt

def plot_decision_regions(X, y, classifier, test_idx=None, resolution=0.02):


    # setup marker generator and color map
    markers = ('s', 'x', 'o', '^', 'v')
    colors = ('red', 'blue', 'lightgreen', 'gray', 'cyan')
    cmap = ListedColormap(colors[:len(np.unique(y))])

    # plot the decision surface
    x1_min, x1_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    x2_min, x2_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx1, xx2 = np.meshgrid(np.arange(x1_min, x1_max, resolution),
                           np.arange(x2_min, x2_max, resolution))
    Z = classifier.predict(np.array([xx1.ravel(), xx2.ravel()]).T)
    Z = Z.reshape(xx1.shape)
    plt.contourf(xx1, xx2, Z, alpha=0.3, cmap=cmap)
    plt.xlim(xx1.min(), xx1.max())
    plt.ylim(xx2.min(), xx2.max())

    for idx, cl in enumerate(np.unique(y)):
        plt.scatter(x=X[y == cl, 0],
                    y=X[y == cl, 1],
                    alpha=0.8,
                    c=colors[idx],
                    marker=markers[idx],
                    label=cl,
                    edgecolor='black')

    # highlight test examples
    if test_idx:
        # plot all examples
        X_test, y_test = X[test_idx, :], y[test_idx]
        plt.scatter(X_test[:, 0],
                    X_test[:, 1],
                    c='',
                    edgecolor='black',
                    alpha=1.0,
                    linewidth=1,
                    marker='o',
                    s=100,
                    label='test set')

```



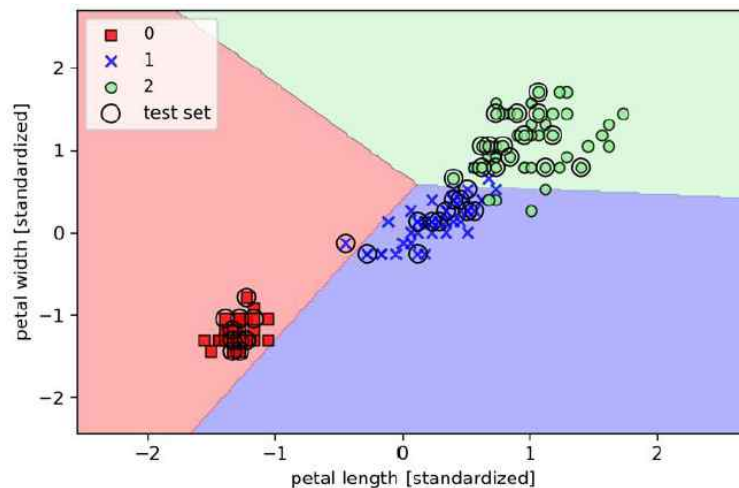
Training a perceptron model using the standardized training data:

```
X_combined_std = np.vstack((X_train_std, X_test_std))
y_combined = np.hstack((y_train, y_test))

plot_decision_regions(X=X_combined_std, y=y_combined,
                      classifier=ppn, test_idx=range(105, 150))
plt.xlabel('petal length [standardized]')
plt.ylabel('petal width [standardized]')
plt.legend(loc='upper left')

plt.tight_layout()
#plt.savefig('images/03_01.png', dpi=300)
plt.show()
```

```
# highlight test examples
if test_idx:
    # plot all examples
    X_test, y_test = X[test_idx, :], y[test_idx]
    plt.scatter(X_test[:, 0],
                X_test[:, 1],
                c='',
                edgecolor='black',
                alpha=1.0,
                linewidth=1,
                marker='o',
                s=100,
                label='test set')
```



- ❑ Three flower classes cannot be perfectly separated by a linear decision boundary.

```
from matplotlib.colors import ListedColorMap
import matplotlib.pyplot as plt

def plot_decision_regions(X, y, classifier, test_idx=None, resolution=0.02):
    # setup marker generator and color map
    markers = ('s', 'x', 'o', '+', '^')
    colors = ('red', 'blue', 'lightgreen', 'gray', 'cyan')
    cmap = ListedColorMap(colors[1:len(np.unique(y))])

    # plot the decision surface
    x1_min, x1_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    x2_min, x2_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx1, xx2 = np.meshgrid(np.arange(x1_min, x1_max, resolution),
                           np.arange(x2_min, x2_max, resolution))
    Z = classifier.predict(np.array([xx1.ravel(), xx2.ravel()]).T)
    Z = Z.reshape(xx1.shape)
    plt.contourf(xx1, xx2, Z, alpha=0.3, cmap=cmap)
    plt.xlim(xx1.min(), xx1.max())
    plt.ylim(xx2.min(), xx2.max())

    for idx, cl in enumerate(np.unique(y)):
        plt.scatter(x=X[y == cl, 0],
                    y=X[y == cl, 1],
                    alpha=0.8,
                    c=colors[idx],
                    marker=markers[idx],
                    label=cl,
                    edgecolor='black')

    # highlight test examples
    if test_idx:
        # plot all examples
        X_test, y_test = X[test_idx, :], y[test_idx]
        plt.scatter(X_test[:, 0],
                    X_test[:, 1],
                    c='',
                    edgecolor='black',
                    alpha=1.0,
                    linewidth=1,
                    marker='o',
                    s=100,
                    label='test set')
```


Logistic Regression

- ❑ a classification model that is very easy to implement and performs very well on linearly separable classes
- ❑ widely used
- ❑ idea behind logistic regression
 - **odds** in favor of a particular event
 - $\text{odds} = p/(1 - p)$ where p = prob. of an event that we want to predict
 - *logit* function: domain ~ 0 to 1 , range \sim entire real-number

$p(y = 1|\mathbf{x})$: conditional prob. that a particular example belongs to class 1 given its features, \mathbf{x} .

- predicts the prob., p , that a certain example belongs to a particular class.

$$\mathbf{Z} = \mathbf{w}^T \mathbf{x} = w_0 x_0 + w_1 x_1 + \cdots + w_m x_m \quad \text{ch 3}$$

Logistic Regression

□ From

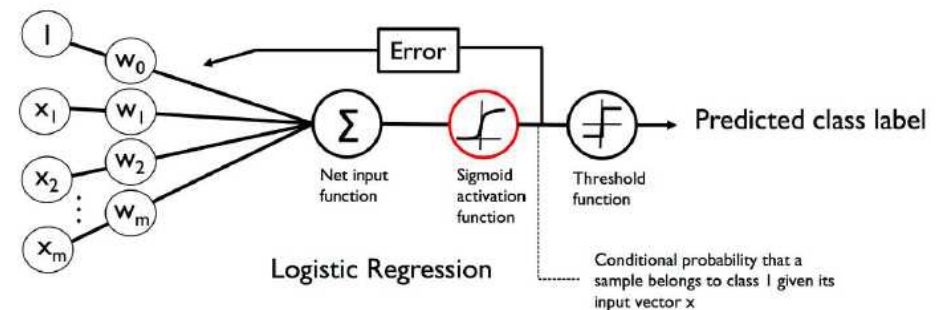

$$\text{logit}(p(y = 1|\mathbf{x})) = w_0x_0 + w_1x_1 + \dots + w_mx_m = \sum_{i=0}^m w_ix_i = \mathbf{w}^T \mathbf{x}$$

□

□

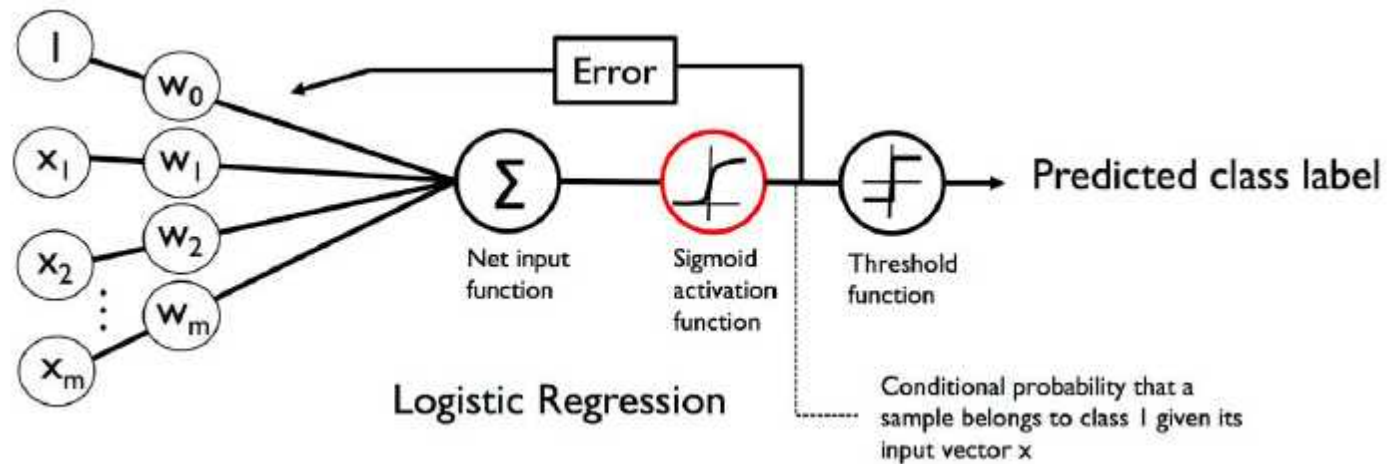
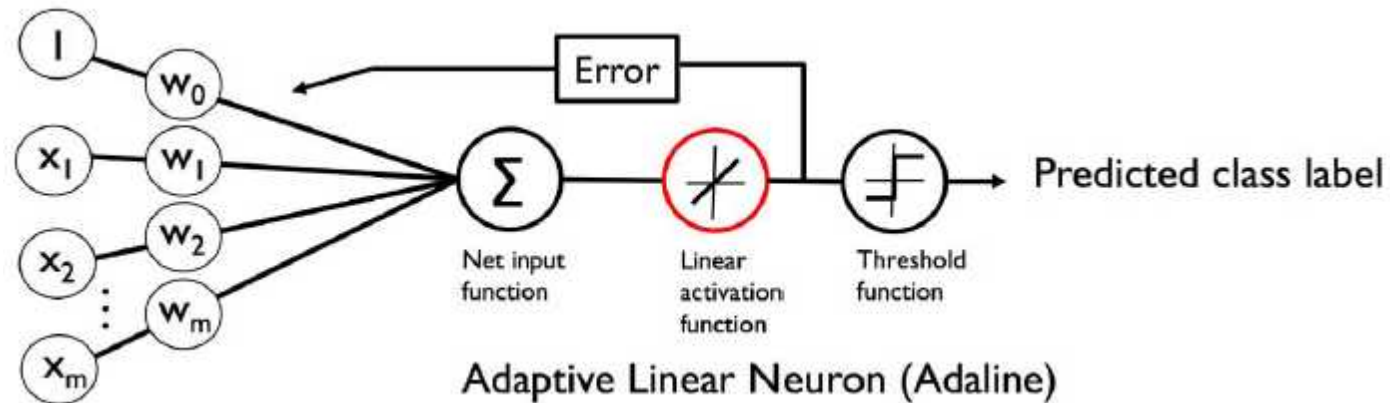
where $Z = \mathbf{w}^T \mathbf{x} = w_0x_0 + w_1x_1 + \dots + w_mx_m$

linear combination of weights and features
 $w_0 \sim \text{bias}, x_0 = 1$





The output of sigmoid function is interpreted as the prob that a particular example belongs to class 1,
 $\phi(z) = p(y = 1|\mathbf{x}; \mathbf{w})$.

Logistic Regression



Logistic Regression

- The predicted probability is converted to a binary outcome



$$p = \phi(z) = \frac{1}{1 + e^{-z}}$$

Updating weights

$$J(\mathbf{w}) = \sum_i \frac{1}{2} (\phi(z^{(i)}) - y^{(i)})^2$$

□ Likelihood

$$L(\mathbf{w}) = P(\mathbf{y} | \mathbf{x}; \mathbf{w}) =$$

Notice

Recall:

The output of sigmoid function is interpreted as the prob that a particular example belongs to class 1, $\phi(z) = p(y = 1 | \mathbf{x}; \mathbf{w})$.

□ log-likelihood

$$\ell(\mathbf{w}) =$$

to be maximized!

□ cost function, $J(\mathbf{w}) =$

$$J(\mathbf{w}) = - \sum_i y^{(i)} \log(\phi(z^{(i)})) + (1 - y^{(i)}) \log(1 - \phi(z^{(i)}))$$

to be minimized!

ch 3

34

Updating weights

partial derivative of log-likelihood

- $$\frac{\partial}{\partial w_j} \ell(w) = \frac{\partial}{\partial w_j} \sum_{i=1}^n \left[y^{(i)} \log(\phi(z^{(i)})) + (1 - \phi(z^{(i)}))^{1-y^{(i)}} \right]$$

$$= \sum_{i=1}^n \frac{\partial}{\partial w_j} \left[y^{(i)} \log(\phi(z^{(i)})) + (1 - \phi(z^{(i)}))^{1-y^{(i)}} \right]$$

- $$\frac{\partial}{\partial w_j} \left[y^{(i)} \log(\phi(z^{(i)})) + (1 - \phi(z^{(i)}))^{1-y^{(i)}} \right]$$

$$= \left(y^{(i)} \frac{1}{\phi(z^{(i)})} - (1 - y^{(i)}) \frac{1}{1 - \phi(z^{(i)})} \right) \frac{\partial}{\partial w_j} \phi(z^{(i)})$$

removing superscript (i)

$$= \left(y \frac{1}{\phi(z)} - (1 - y) \frac{1}{1 - \phi(z)} \right) \frac{\partial}{\partial w_j} \phi(z)$$

$$= \left(y \frac{1}{\phi(z)} - (1 - y) \frac{1}{1 - \phi(z)} \right) \phi(z)(1 - \phi(z)) \frac{\partial}{\partial w_j} z$$

$$= \left(y(1 - \phi(z)) - (1 - y)\phi(z) \right) x_j$$

$$= \underline{(y - \phi(z))x_j}$$

- $$w_j := w_j + \eta \sum_{i=1}^n (y^{(i)} - \phi(z^{(i)})) x_j^{(i)}$$

Updating weights

- maximizing log-likelihood is equal to minimizing the cost function $J(w)$
- gradient descent update rule

$$\Delta w_j =$$

Converting Adaline implementation to algorithm for logistic regression

```
class LogisticRegressionGD(object):
    """Logistic Regression Classifier using gradient descent.

    Parameters
    -----
    eta : float
        Learning rate (between 0.0 and 1.0)
    n_iter : int
        Passes over the training dataset.
    random_state : int
        Random number generator seed for random weight
        initialization.

    Attributes
    -----
    w_ : 1d-array
        Weights after fitting.
    cost_ : list
        Logistic cost function value in each epoch.

    """
```

```
def net_input(self, X):
    """Calculate net input"""
    return np.dot(X, self.w_[1:]) + self.w_[0]

def activation(self, z):
    """Compute logistic sigmoid activation"""
    return 1. / (1. + np.exp(-np.clip(z, -250, 250)))

def predict(self, X):
    """Return class label after unit step"""
    return np.where(self.net_input(X) >= 0.0, 1, 0)
    # equivalent to:
    # return np.where(self.activation(self.net_input(X)) >= 0.5, 1, 0)
```



```

def __init__(self, eta=0.05, n_iter=100, random_state=1):
    self.eta = eta
    self.n_iter = n_iter
    self.random_state = random_state

def fit(self, X, y):
    """ Fit training data.

    Parameters
    -----
    X : {array-like}, shape = [n_examples, n_features]
        Training vectors, where n_examples is the number of examples and
        n_features is the number of features.
    y : array-like, shape = [n_examples]
        Target values.

    Returns
    -----
    self : object

    """
    rgen = np.random.RandomState(self.random_state)
    self.w_ = rgen.normal(loc=0.0, scale=0.01, size=1 + X.shape[1])
    self.cost_ = []

    for i in range(self.n_iter):
        net_input = self.net_input(X)
        output = self.activation(net_input)
        errors = (y - output)
        self.w_[1:] += self.eta * X.T.dot(errors)
        self.w_[0] += self.eta * errors.sum()

        # note that we compute the logistic 'cost' now
        # instead of the sum of squared errors cost
        cost = -y.dot(np.log(output)) - ((1 - y).dot(np.log(1 - output)))
        self.cost_.append(cost)

    return self

```

Logistic regression

- ❑ Consider only Iris-setosa and Iris-versicolor flowers (classes 0 and 1)

```
X_train_01_subset = X_train[(y_train == 0) | (y_train == 1)]
y_train_01_subset = y_train[(y_train == 0) | (y_train == 1)]

lrgd = LogisticRegressionGD(eta=0.05, n_iter=1000, random_state=1)
lrgd.fit(X_train_01_subset,
        y_train_01_subset)

plot_decision_regions(X=X_train_01_subset,
                     y=y_train_01_subset,
                     classifier=lrgd)

plt.xlabel('petal length [standardized]')
plt.ylabel('petal width [standardized]')
plt.legend(loc='upper left')

plt.tight_layout()
#plt.savefig('images/03_05.png', dpi=300)
plt.show()
```

