

Training a logistic regression model with scikit-learn

- ❑ how to use scikit-learn's more optimized implementation of logistic regression, which also supports multiclass settings off the shelf.
 - optimization algorithms - newton-cg, lbfgs, liblinear, sag, and saga, in addition to SGD
 - lbfgs - limited-memory Broyden-Fletcher-Goldfarb-Shanno algorithm
 - input parameter: solver = 'lbfgs'
 - multiclass classification - multinomial or OvR can be chosen
 - input parameter: multi_class = 'ovr', or 'multinomial'

Training a logistic regression model with scikit-learn

❑ tackling overfitting via regularization

- introduce additional information (bias) to penalize extreme parameter (weight) values.
- L2 regularization (shrinkage or weight decay)

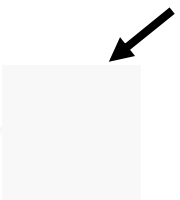
$$\frac{\lambda}{2} \|\mathbf{w}\|^2 = \frac{\lambda}{2} \sum_{j=1}^m w_j^2$$

λ - regularization parameter

- For regularization to work properly, ensure that all features are on comparable scales.

$$J(\mathbf{w}) = \sum_{i=1}^n \left[-y^{(i)} \log(\phi(z^{(i)})) - (1 - y^{(i)}) \log(1 - \phi(z^{(i)})) \right] +$$


regularization
strength



- via the regularization parameter λ , we can control how well we fit the training data, while keeping the weights small

support vector machines


- SVMs use a boundary called a hyperplane to partition data into two groups of similar class values.



case: samples are linearly separable.


Hyperplane and halfspace

- hyperplane H in E^n is a set of x such that $px = k$ where p is a nonzero vector in E^n and normal to the hyperplane, and k is scalar




Hyperplane and halfspace

- hyperplane divides E^n into two regions, called halfspaces.
- halfspace is a collection of points of the form $\{x: px \geq k\}$
- another halfspace is a collection of points of the form $\{x: Px \leq k\}$



support vector machines

- In two dimensions, the task of the SVM algorithm is to identify a line that separates the two classes.




case: samples are linearly separable.

- Question is how does the algorithm choose the most appropriate one.

Answer - find the maximum margin hyperplane

support vectors and MMH

- ❑ maximum margin hyperplane (MMH) creates the greatest separation between the two classes.
- ❑ MMH will generalize best to future data.
- ❑ support vectors are the points from each class that are the closest to the MMH.
- ❑ support vectors alone define the MMH.
 - support vectors provide a very compact way to store a classification model.




How to determine the MMH?

- general equation of a hyperplane in \mathbb{R}^N

- scale w and b such that $\min |wx + b| = 1$ over all training data.
 - The corresponding hyperplane is called canonical hyperplane.

$$w \cdot x + b = +1 \text{ or } w \cdot x + b = -1$$

- For a canonical hyperplane, margin ρ is given by $1 / \|w\|$



How to determine the MMH?


□ case: classes are linearly separable

hyperplane: $w = \{w_1, w_2, \dots, w_n\}$
 b : intercept, scalar

$$\vec{w} \cdot \vec{x} + b = 0$$

For a canonical hyperplane, we have

Thus $x^{(i)}$ is correctly classified when




Find a hyperplane satisfying these with the biggest margin.

How to determine the MMH?

case: nonlinearly separable data


- What happens in the case that the data are not linearly separable?
- the constraints imposed in the linearly separable case cannot be met.



How to determine the MMH?

case: nonlinearly separable data

- Solution : creates a soft margin that allows some points to fall on the incorrect side of the margin




How to determine the MMH?

case: nonlinearly separable data


□ problem formulation

$$\begin{aligned} \min \quad & \frac{1}{2} \|\vec{w}\|^2 + C \sum_{i=1}^n \xi_i \\ \text{s.t.} \quad & y_i(\vec{w} \cdot \vec{x}_i + b) \geq 1 - \xi_i, \forall \xi_i \geq 0 \end{aligned}$$



Kernel methods for linearly inseparable data

- In many real-world datasets, relationship between variables are nonlinear. Using some mapping technique, a nonlinear relationship can be made quite linear.
- example dataset not linearly separable, XOR data




Kernel methods for linearly inseparable data

- ❑ transform a linearly inseparable data to the data linearly separable
- ❑ uses nonlinear combinations of the original features
- ❑ original space is transformed to a higher-dimensional space via a mapping function, ϕ
- ❑ **example** - two dimensional dataset is transformed into a new three-dimensional feature space, where the classes become separable


$$\phi(x_1, x_2) = (z_1, z_2, z_3) = (x_1, x_2, x_1^2 + x_2^2)$$

Kernel methods for linearly inseparable data



$$\phi(x_1, x_2) = (z_1, z_2, z_3) = (x_1, x_2, x_1^2 + x_2^2)$$

nonlinear function of
the features x_1 and x_2



Kernel SVM using kernel tricks to find separating hyperplanes in a higher-dimensional space

- to save the expensive step of calculating the dot product $\phi(\mathbf{x}^{(i)})^T \phi(\mathbf{x}^{(j)})$ between two points explicitly, kernel function is defined.

$$\underline{\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)})} \triangleq \phi(\mathbf{x}^{(i)})^T \phi(\mathbf{x}^{(j)})$$

- radial basis function (RBF) kernel
 - widely used, a.k.a Gaussian kernel

- often



$$\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \exp\left(-\gamma \|\mathbf{x}^{(i)} - \mathbf{x}^{(j)}\|^2\right)$$

- γ a free parameter to be optimized

Kernel SVM using kernel tricks to find separating hyperplanes in a higher-dimensional space

```
svm = SVC(kernel='rbf', random_state=1, gamma=0.10, C=10.0)
svm.fit(X_xor, y_xor)
plot_decision_regions(X_xor, y_xor,
                     classifier=svm)

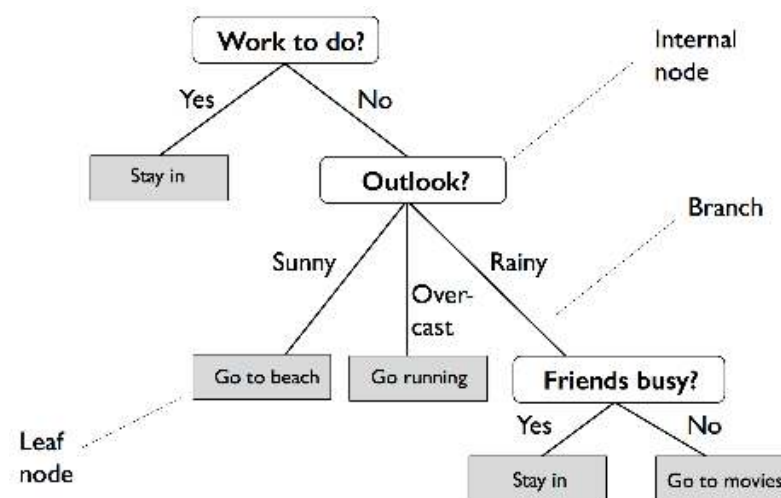
plt.legend(loc='upper left')
plt.tight_layout()
#plt.savefig('images/03_14.png', dpi=300)
plt.show()
```



XOR data

Decision tree learning

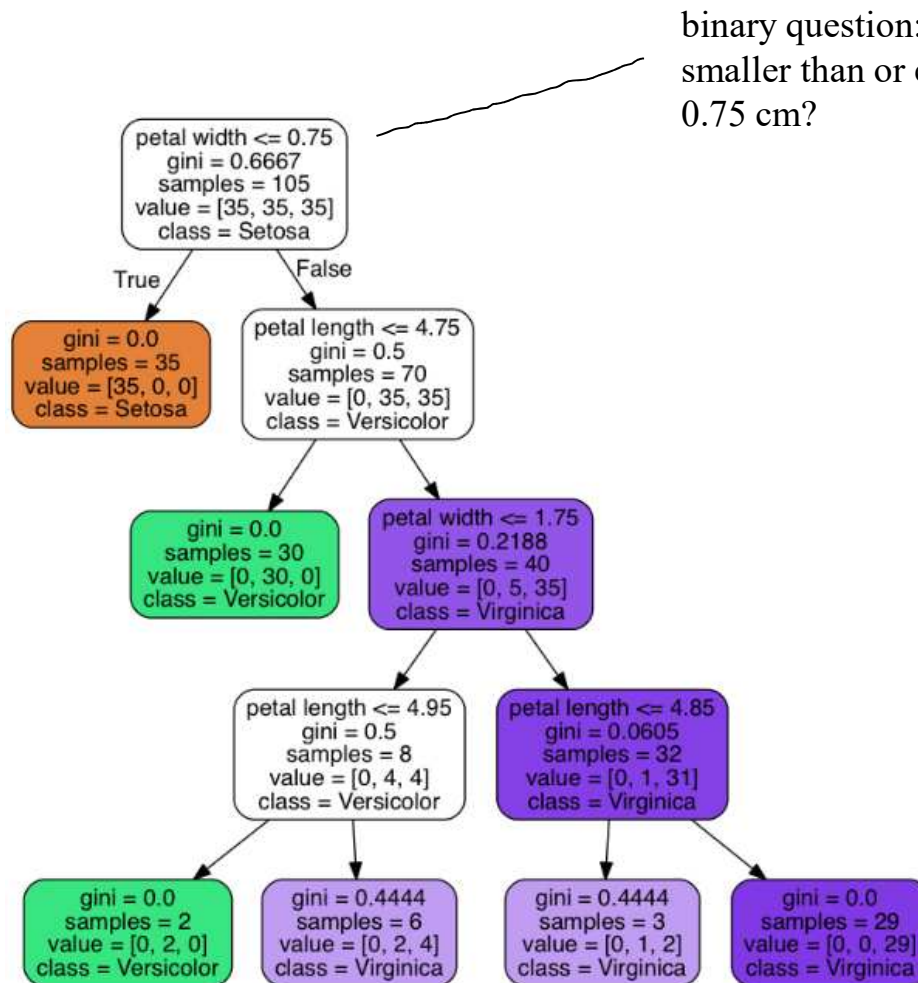
- Decision tree classifiers are attractive model if we are about interpretability



- **Question:** which feature to split upon?
 - degree to which a subset of examples contains only a single class is known as **purity**.

p. 132, machine learning with R, 3rd ed.


Decision tree learning



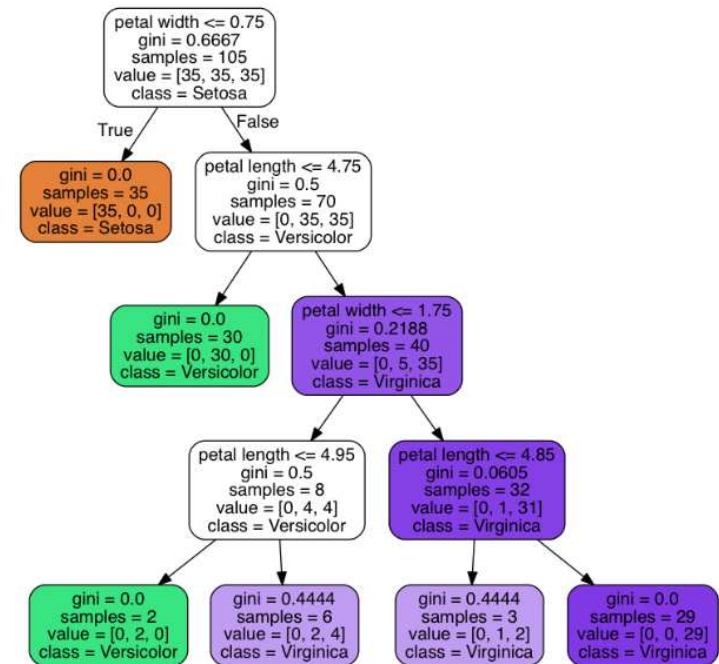
- splits data on the feature that results in the largest information gain (IG)
- this can result in a very deep tree with many nodes
 - this can lead to overfitting
 - typically want to prune the tree by setting a limit for the maximal depth of the tree.

maximizing information gain

- ❑ objective: to split the nodes at the most informative features
- ❑ objective function
 - maximizes the IG at each split



- f : feature to perform the split
- D_p : data set of parent
- D_j : data set of j -th child node
- $I()$: **impurity** measure
- N_p : total number of training examples at parent node
- N_j : number of examples in j -th child node




maximizing information gain

- The lower the impurities of the child nodes, the larger the information gain

$$IG(D_p, f) = I(D_p) - \sum_{j=1}^m \frac{N_j}{N_p} I(D_j)$$

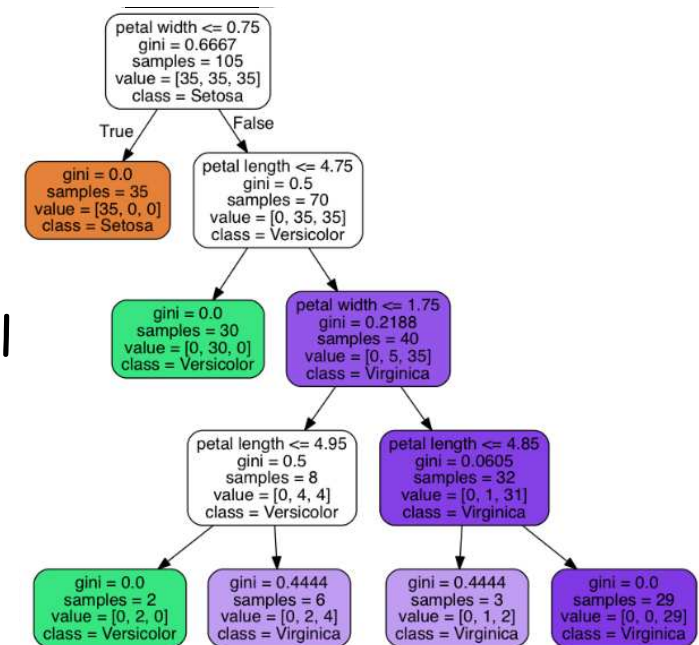
- for simplicity and to reduce combinatorial search space, most libraries implement *binary* decision trees.

- each parent node is split into two child nodes, D_{left} and D_{right}




- common impurity measures

- Gini impurity, entropy, classification error



Impurity measures

- $p(i|t)$ ~ proportion of the examples that belong to the class i for a particular node t .
- **entropy** for all non-empty classes (c : number of classes)



- entropy is zero if all examples at a node belong to the same class
- entropy is maximal if we have a uniform class distribution
- entropy criterion attempts to maximize the mutual information in the tree (**minimize entropy**)

Impurity measures

- **Gini impurity** : criterion to minimize the probability of misclassification

- similar to entropy, Gini impurity is maximal if the classes are perfectly mixed.


in a binary class setting ($c = 2$):

$$I_G(t) = 1 - \sum_{i=1}^c 0.5^2 = 0.5$$

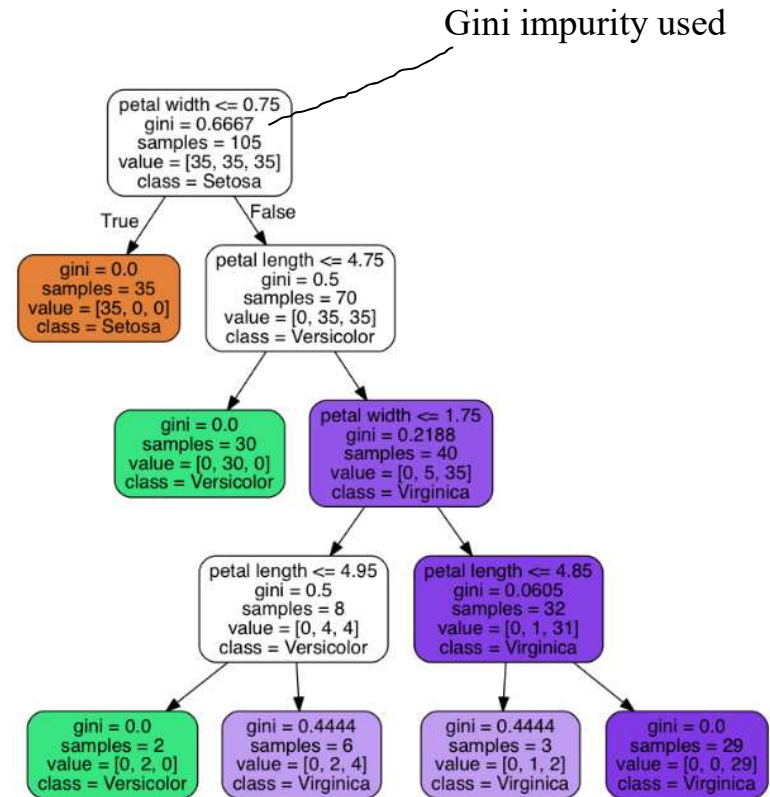
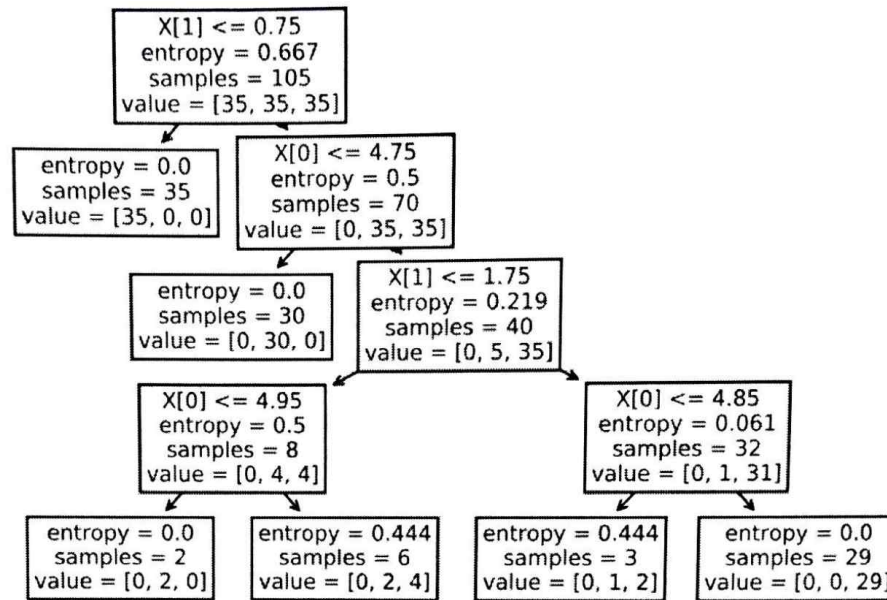
- **classification error** : useful criterion for pruning but not recommended for growing a decision tree, (less sensitive to changes in the class probabilities of the nodes)

$$I_E(t) = 1 - \max\{p(i|t)\}$$

decision boundaries of decision tree model




Visualization of the decision tree model



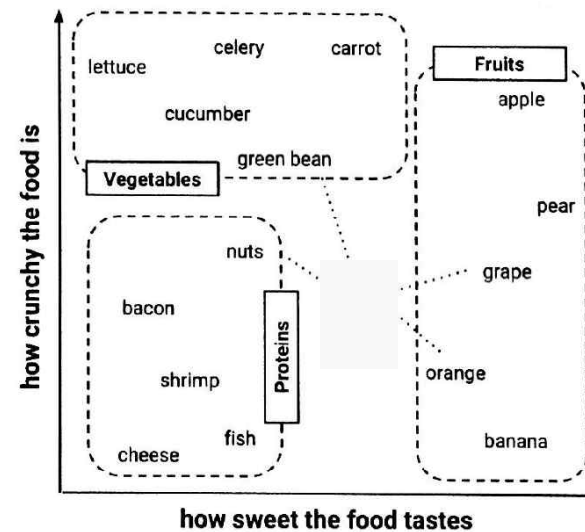
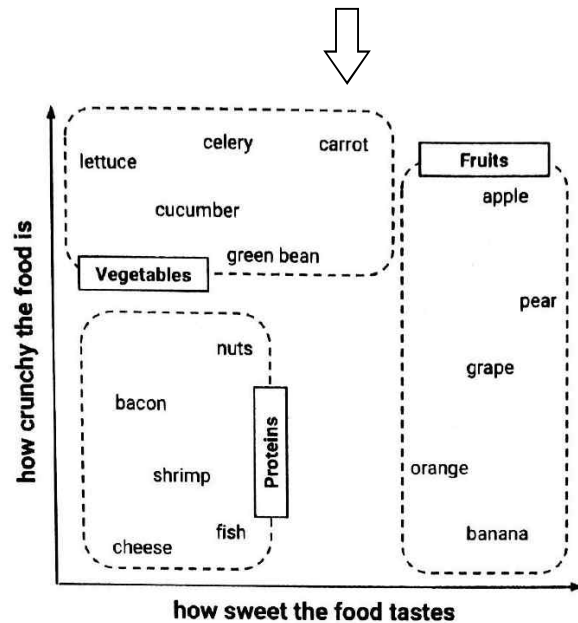
K-nearest neighbors classifier

- ❑ KNN
 - ❑ a typical example of a lazy learner
 - ❑ decision steps
 - choose the number k and a distance metric
 - find the k -nearest neighbors of the training data
 - assign the class label by majority voting
- ❑ following figure illustrates how a new data point(?) is assigned the triangle class label based on majority voting among its five nearest neighbors.



K-nearest neighbors classifier

Ingredient	Sweetness	Crunchiness	Food type
Apple	10	9	Fruit
Bacon	1	4	Protein
Banana	10	1	Fruit
Carrot	7	10	Vegetable
Celery	3	10	Vegetable
Cheese	1	1	Protein



KNN model in scikit-learning

```
from sklearn.neighbors import KNeighborsClassifier

knn = KNeighborsClassifier(n_neighbors=
                           p=
                           metric=
knn.fit(X_train_std, y_train)

plot_decision_regions(X_combined_std, y_combined,
                      classifier=knn, test_idx=range(105, 150))

plt.xlabel('petal length [standardized]')
plt.ylabel('petal width [standardized]')
plt.legend(loc='upper left')
plt.tight_layout()
#plt.savefig('images/03_24.png', dpi=300)
plt.show()
```

